

# SCALING ERLANG WEB APPLICATIONS

## 100 TO 100K USERS AT ONE WEB SERVER

Fernando Benavides (*@elbrujothalcon*)

Inaka Labs

March 6, 2012

# HELLO WORLD!

- I'm a developer since I was 10
- I'm an Erlang developer since 2008
- I've worked in many dynamic web sites
- Most of them with high scale requirements
- I'll share my experience with you

# HELLO WORLD!

- I'm a developer since I was 10
- I'm an Erlang developer since 2008
- I've worked in many dynamic web sites
- Most of them with high scale requirements
- I'll share my experience with you

# HELLO WORLD!

- I'm a developer since I was 10
- I'm an Erlang developer since 2008
- I've worked in many dynamic web sites
- Most of them with high scale requirements
- I'll share my experience with you

# HELLO WORLD!

- I'm a developer since I was 10
- I'm an Erlang developer since 2008
- I've worked in many dynamic web sites
- Most of them with high scale requirements
- I'll share my experience with you

# HELLO WORLD!

- I'm a developer since I was 10
- I'm an Erlang developer since 2008
- I've worked in many dynamic web sites
- Most of them with high scale requirements
- I'll share my experience with you

# OUTLINE

## THE CHALLENGE

What do we have to deal with?

## THE PLAN

How do we face it?

## THE TIPS AND TRICKS

What have we learned from it?

# OUTLINE

## THE CHALLENGE

What do we have to deal with?

## THE PLAN

How do we face it?

## THE TIPS AND TRICKS

What have we learned from it?



# OUTLINE

## THE CHALLENGE

What do we have to deal with?

## THE PLAN

How do we face it?

## THE TIPS AND TRICKS

What have we learned from it?

We will work on the scalability of a *web* project that has an *HTTP API* and keeps clients *connected* to the server for *long periods* of time.

Examples:

- Social sites
- Chat sites
- Sports sites

We will work on the scalability of a *web* project that has an *HTTP API* and keeps clients *connected* to the server for *long periods* of time.

Examples:

- Social sites
- Chat sites
- Sports sites

We will work on the scalability of a *web* project that has an *HTTP API* and keeps clients *connected* to the server *for long periods* of time.

Examples:

- Social sites
- Chat sites
- Sports sites

We will work on the scalability of a *web* project that has an *HTTP API* and keeps clients *connected* to the server for *long periods* of time.

Examples:

- Social sites
- Chat sites
- Sports sites

We will work on the scalability of a *web* project that has an *HTTP API* and keeps clients *connected* to the server for *long periods* of time.

Examples:

- Social sites
- Chat sites
- Sports sites

*We will focus on*

- OTP behaviours
- TCP connections
- mochiweb
- Underlying system configurations

*We will **not** deal with*

- Multiple machines/nodes
- Databases

*We will focus on*

- OTP behaviours
- TCP connections
- mochiweb
- Underlying system configurations

*We will **not** deal with*

- Multiple machines/nodes
- Databases



# THE PLAN

# GENERAL CONSIDERATIONS

- Create a system that **works**
- Automate your clients
- Keep a human watching
- Be patient

# GENERAL CONSIDERATIONS

- Create a system that **works**
- **Automate your clients**
- Keep a human watching
- Be patient

# GENERAL CONSIDERATIONS

- Create a system that **works**
- Automate your clients
- Keep a human watching
- Be patient

# GENERAL CONSIDERATIONS

- Create a system that **works**
- Automate your clients
- Keep a human watching
- Be patient

# GENERAL CONSIDERATIONS

- Create a system that **works**
- Automate your clients
- Keep a human watching
- Be patient

# GOALS

- Test the system as it is
- How many users can the system handle **as is**?
- Find  $N$  and  $C$

# STEPS

- Choose  $N$  and  $C$
- Test the API
- Test long-lived connections
- Test both
- Repeat with higher values for  $N$  and  $C$



# STEPS

- Choose  $N$  and  $C$
- Test the API
- Test long-lived connections
- Test both
- Repeat with higher values for  $N$  and  $C$

# GOALS

- Improve the system environment
- Tune-In the machine(s)
- **Don't** touch the code

# STEPS

- Check kernel variables
- Check system limits
- Check Erlang VM parameters

# GOALS

- Tune up **your** system
- Discover scalability issues and fix them
- Find the biggest  $N$  and  $C$  for **one node**

# STEPS

- Choose  $N$  and  $C$  to fail
- Find a problem
- Fix it
- Add it to the list of *Tips and Tricks*
- Repeat with higher values for  $N$  and  $C$

# STEPS

- Choose  $N$  and  $C$  to fail
- Find a problem
- Fix it
- Add it to the list of *Tips and Tricks*
- Repeat with higher values for  $N$  and  $C$

# GOALS

- Get the system ready to work on many nodes
- Design the system topology
- Find  $N$  and  $C$  **per node**

# STEPS

- Get the second node running
- Choose  $N$  and  $C$
- Try interconnected instances
- Try independent instances
- Repeat with higher values for  $N$  and  $C$



# STEPS

- Get the second node running
- Choose  $N$  and  $C$
- Try interconnected instances
- Try independent instances
- Repeat with higher values for  $N$  and  $C$

# OS TWEAKS

## *Kernel Variables*

```
sysctl -w net.ipv4.ip_local_port_range="1024 65535"  
sysctl -w net.core.rmem_max=16777216  
sysctl -w net.core.wmem_max=16777216  
sysctl -w net.ipv4.tcp_rmem="4096 87380 16777216"  
sysctl -w net.ipv4.tcp_wmem="4096 65536 16777216"  
sysctl -w net.ipv4.tcp_syncookies=1  
sysctl -w net.ipv4.tcp_mem="50576 64768 98152"  
sysctl -w net.core.netdev_max_backlog=2500  
sysctl -w net.netfilter.nf_conntrack_max=1233000
```

## *Open Files Limit*

```
ulimit -n 999999
```

## *Erlang VM tweaks*

- +P Number of Processes
- +K Kernell Polling
- SMP SMP Support

# OS TWEAKS

## *Kernel Variables*

```
sysctl -w net.ipv4.ip_local_port_range="1024 65535"  
sysctl -w net.core.rmem_max=16777216  
sysctl -w net.core.wmem_max=16777216  
sysctl -w net.ipv4.tcp_rmem="4096 87380 16777216"  
sysctl -w net.ipv4.tcp_wmem="4096 65536 16777216"  
sysctl -w net.ipv4.tcp_syncookies=1  
sysctl -w net.ipv4.tcp_mem="50576 64768 98152"  
sysctl -w net.core.netdev_max_backlog=2500  
sysctl -w net.netfilter.nf_conntrack_max=1233000
```

## *Open Files Limit*

```
ulimit -n 999999
```

## *Erlang VM tweaks*

- +P Number of Processes
- +K Kernell Polling
- SMP SMP Support

# OS TWEAKS

## *Kernel Variables*

```
sysctl -w net.ipv4.ip_local_port_range="1024 65535"  
sysctl -w net.core.rmem_max=16777216  
sysctl -w net.core.wmem_max=16777216  
sysctl -w net.ipv4.tcp_rmem="4096 87380 16777216"  
sysctl -w net.ipv4.tcp_wmem="4096 65536 16777216"  
sysctl -w net.ipv4.tcp_syncookies=1  
sysctl -w net.ipv4.tcp_mem="50576 64768 98152"  
sysctl -w net.core.netdev_max_backlog=2500  
sysctl -w net.netfilter.nf_conntrack_max=1233000
```

## *Open Files Limit*

```
ulimit -n 999999
```

## *Erlang VM tweaks*

- +P Number of Processes
- +K Kernell Polling
- SMP SMP Support

# ERLANG TWEAKS

TODO: Copy from the article on listeners  
TODO: Copy from the article on inbound TCP connections  
TODO: Copy from the article on outbound TCP connections

## GEN\_EVENT

TODO: Copy from the article on sup\_handler  
TODO: Copy from the article on long delivery queues

## GEN\_SERVERS

TODO: Copy from the article on timing out  
TODO: Copy from the article on too much memory  
TODO: Copy from the article on taking too long to initialize

# SUPERVISORS

TODO: Copy from the article



# PROCESS REGISTRATION

TODO: Copy from the article

# TIMERS

TODO: Copy from the article

# LOGGING

TODO: Copy from the article

# SUMMARY

TODO: Summary

## OTHER STUFF

THAT WE LEFT OUT OF THIS PRESENTATION

TODO: List of other scalability stuff we left out

Any questions?

```
-spec fact(integer()) -> integer().  
fact(N) ->  
    lists:fold(fun(X, F) ->  
                F * X  
            end, 1, lists:seq(1,N)).
```