

# SCALING ERLANG WEB APPLICATIONS

## 100 TO 100K USERS AT ONE WEB SERVER

Fernando Benavides (*@elbrujohalcon*)

Inaka Labs

March 26, 2012



# HELLO WORLD!

- I'm a developer since I was 10
- I worked with Visual Basic, C#, .NET, Javascript ...
- I switched to functional programming in 2008
- I wrote my thesis project in Haskell
- I'm an Erlang developer since then



# HELLO WORLD!

- I'm a developer since I was 10
- I worked with Visual Basic, C#, .NET, Javascript . . .
- I switched to functional programming in 2008
- I wrote my thesis project in Haskell
- I'm an Erlang developer since then



# HELLO WORLD!

- I'm a developer since I was 10
- I worked with Visual Basic, C#, .NET, Javascript . . .
- I switched to functional programming in 2008
- I wrote my thesis project in Haskell
- I'm an Erlang developer since then



# HELLO WORLD!

- I'm a developer since I was 10
- I worked with Visual Basic, C#, .NET, Javascript ...
- I switched to functional programming in 2008
- I wrote my thesis project in Haskell
- I'm an Erlang developer since then



# HELLO WORLD!

- I'm a developer since I was 10
- I worked with Visual Basic, C#, .NET, Javascript . . .
- I switched to functional programming in 2008
- I wrote my thesis project in Haskell
- I'm an Erlang developer since then



# INAKA



# INTRODUCTION

My talk is on the scalability of a *web* project that has an *HTTP API* and a component that keeps clients *connected* to the server for *long periods* of time.

It's a design pattern seen in many places:

- Chat Applications
- Social Sites
- Sport Sites





# INTRODUCTION

My talk is on the scalability of a *web* project that has an *HTTP API* and a component that keeps clients *connected* to the server for *long periods* of time.

It's a design pattern seen in many places:

- Chat Applications
- Social Sites
- Sport Sites



# INTRODUCTION

My talk is on the scalability of a *web* project that has an *HTTP API* and a component that keeps clients *connected* to the server for *long periods* of time.

It's a design pattern seen in many places:

- Chat Applications
- Social Sites
- Sport Sites



# INTRODUCTION

My talk is on the scalability of a *web* project that has an *HTTP API* and a component that keeps clients *connected* to the server for *long periods* of time.

It's a design pattern seen in many places:

- Chat Applications
- Social Sites
- Sport Sites



# SCOPE

*We will improve the way we use*

- OTP behaviours
- TCP and HTTP connections
- Underlying system configurations

*We will **not** deal with*

- Multiple machines/nodes
- Database choices and/or implementations



# SCOPE

*We will improve the way we use*

- OTP behaviours
- TCP and HTTP connections
- Underlying system configurations

*We will **not** deal with*

- Multiple machines/nodes
- Database choices and/or implementations



# MATCH STREAM

## GENERAL IDEA

*A soccer match is played at some stadium*



# MATCH STREAM

## GENERAL IDEA

*Soccer fans are connected to the internet in their offices*



# MATCH STREAM

## GENERAL IDEA

*A reporter is at the stadium with his device*






# MATCH STREAM

## GENERAL IDEA

MATCHSTREAM ***connects*** them in ***real time***



# MATCH STREAM

## REQUIREMENTS

### SYSTEM CHALLENGES

- Many concurrent users connecting at the same time
- Two-hour-long bursts of connections followed by long periods of inactivity
- Real-time updates

Erlang seems to be the right fit for this



# MATCH STREAM

## REQUIREMENTS

### SYSTEM CHALLENGES

- Many concurrent users connecting at the same time
- Two-hour-long bursts of connections followed by long periods of inactivity
- Real-time updates

Erlang seems to be **the right fit for this**



# MATCH STREAM

## REQUIREMENTS

### SYSTEM CHALLENGES

- Many concurrent users connecting at the same time
- Two-hour-long bursts of connections followed by long periods of inactivity
- Real-time updates

Erlang seems to be **the right fit for this**



# MATCH STREAM

## REQUIREMENTS

### SYSTEM CHALLENGES


- Many concurrent users connecting at the same time
- Two-hour-long bursts of connections followed by long periods of inactivity
- Real-time updates

Erlang seems to be **the right fit for this**



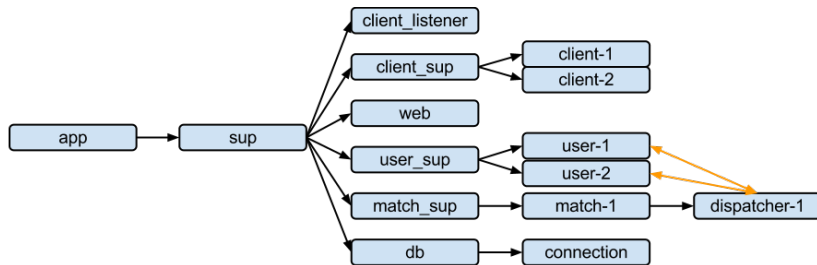
# MATCH STREAM

## GENERAL DESIGN

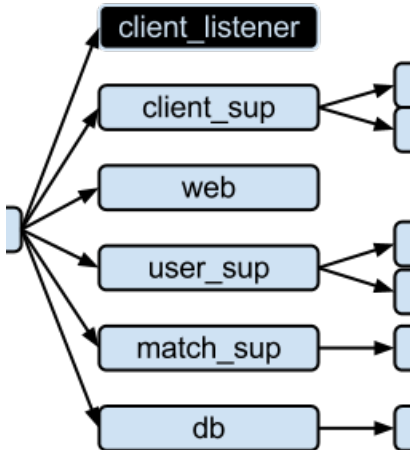


# MATCH STREAM

## ARCHITECTURE



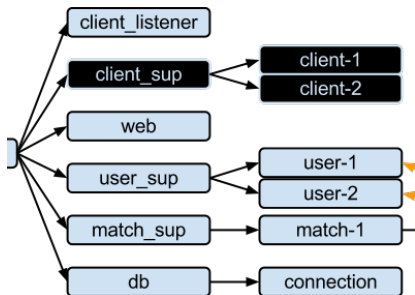
# COMPONENTS



`CLIENT_LISTENER` `gen_server`.  
Listens on a TCP  
port to receive  
client connections



# COMPONENTS

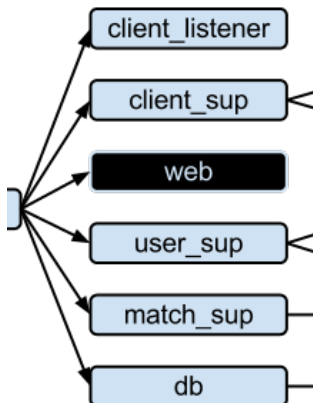


**CLIENT\_SUP** supervisor.  
 Supervises  
 connection  
 processes

**CLIENT** gen\_fsm.  
 Handles a TCP  
 connection




# COMPONENTS



WEB mochiweb server.  
Listens for HTTP  
API calls



# COMPONENTS

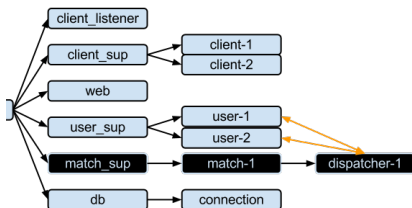


`USER_SUP` supervisor.  
 Supervises user processes

`USER` gen\_server.  
 Subscribes to match dispatchers and sends events to clients



# COMPONENTS




**MATCH\_SUP** supervisor.  
 Supervises match processes

**MATCH** gen\_server.  
 Listens to match events, stores them

**DISPATCHER** gen\_event dispatcher.  
 Delivers match events



# COMPONENTS



DB `gen_server`.  
 Processes  
 database  
 operations

CONNECTION `erldis client`.  
 Handles the  
 connection to the  
 database



## LESSON LEARNED

*Simply using Erlang to build your system is **not enough** to ensure scalability*



# MEASURES

- N** *Connections*. Number of connections the server can handle
- C** *Concurrency*. Number of multiple connections starting at a time
- ART** *Average Response Time*. How much does it take for the server to send an event



# TOOLS

## TEST CLIENT

We create our own test client for TCP connections

## APACHEBENCH

To test API calls

## ENTOP

We use it to see what's going on in the server





# STAGE 0

## ESTABLISHING A BASELINE

### GOALS

- Find how much the system can handle

### STEPS

- Create automated testers
- Start the system on a *clean* machine
- Test repeatedly adjusting the number of connections
- Have a human using the system himself



# STAGE 0

## ESTABLISHING A BASELINE

### GOALS

- Find how much the system can handle


### STEPS

- Create automated testers
- Start the system on a *clean* machine
- Test repeatedly adjusting the number of connections
- Have a human using the system himself



# STAGE 1

## RESULTS



N 1000  
C 5  
ART 26s

## STAGE 1

## TUNE THE OS AND THE VM

## GOALS

- Improve the underlying Operating System
- Improve the Erlang VM Configuration



# STAGE 1

## TUNE THE OS AND THE VM

### GOALS

- Improve the underlying Operating System
- Improve the Erlang VM Configuration

### SETTINGS TO TUNE UP

- Open files limit
- TCP connections limit
- TCP backlog size
- TCP memory allocation
- Number of Erlang processes



# STAGE 1

## TUNE THE OS AND THE VM

### GOALS

- Improve the underlying Operating System
- Improve the Erlang VM Configuration

### SETTINGS TO TUNE UP

- Open files limit
- TCP connections limit
- TCP backlog size
- TCP memory allocation
- Number of Erlang processes



# STAGE 1

## TUNE THE OS AND THE VM

### GOALS

- Improve the underlying Operating System
- Improve the Erlang VM Configuration

### SETTINGS TO TUNE UP

- Open files limit
- TCP connections limit
- TCP backlog size
- TCP memory allocation
- Number of Erlang processes



# STAGE 1

## TUNE THE OS AND THE VM

### GOALS

- Improve the underlying Operating System
- Improve the Erlang VM Configuration

### SETTINGS TO TUNE UP

- Open files limit
- TCP connections limit
- TCP backlog size
- TCP memory allocation
- Number of Erlang processes





# STAGE 1

## TUNE THE OS AND THE VM

### GOALS

- Improve the underlying Operating System
- Improve the Erlang VM Configuration


### SETTINGS TO TUNE UP

- Open files limit
- TCP connections limit
- TCP backlog size
- TCP memory allocation
- Number of Erlang processes



# STAGE 1

## RESULTS



N 4000  
C 5  
ART 35s

## STAGE 2

### IMPROVING MATCH STREAM

We can't blame the machine anymore, we need  
to improve **our system**



## STAGE 2.1

### CONNECTION TWEAKS

#### BACKLOG

- Allow more concurrent connections
- Don't forget TCP tuning your HTTP server



## STAGE 2.1

### CONNECTION TWEAKS

#### CLIENT\_LISTENER

```
gen_tcp:listen(Port,  
  [binary, {packet, line}, {keepalive, true},  
   {active, false}, {reuseaddr, true},  
   {backlog, 128000}, {send_timeout, 32000},  
   {send_timeout_close, true}]).
```

#### WEB

```
mochiweb_http:start(  
  [{name, ?MODULE}, {loop, {?MODULE, loop}},  
   {backlog, 128000}, {port, Port}]).
```



# STAGE 2.1

## CONNECTION TWEAKS

### OUTBOUND CONNECTIONS

- For instance, database connections
- Don't use just one of them
- You may have separated connections for different purposes



# STAGE 2.1

## CONNECTION TWEAKS

```
-define(REDIS_CONNECTIONS, 200).  
-record(state, {redis :: [pid()] }).  
  
...  
Redis =  
    lists:map(  
        fun(_) ->  
            {ok, Conn} = erldis_client:start_link()  
            Conn  
        end, lists:seq(1, ?REDIS_CONNECTIONS) ),  
{ok, #state{redis = Redis}}.
```



## STAGE 2.1

### CONNECTION TWEAKS

```
handle_call(Request, From, State) ->
  [RedisConn|Redis] = State#state.redis,
  proc_lib:spawn_link(
    fun() ->
      Res = handle_call(Request, RedisConn),
      gen_server:reply(From, Res)
    end),
  {noreply, State#state{redis =
    Redis ++ [RedisConn] }}.
```





## STAGE 2.1

### CONNECTION TWEAKS

#### LISTENERS

- You can listen to more than one port
- For unified urls, use *nginx* in front of the server



## STAGE 2.1

### CONNECTION TWEAKS

```
init([]) ->
...
    Listeners =
        [{list_to_atom("client-listener-" ++
                        integer_to_list(I)),
         client_listener, start_link, [I],
         permanent, brutal_kill, worker,
         [client_listener]}
         || I <- lists:seq(MinPort, MaxPort)],
        {ok, {{one_for_one, 5, 10}, Listeners}}.
```



## STAGE 2.1

### CONNECTION TWEAKS

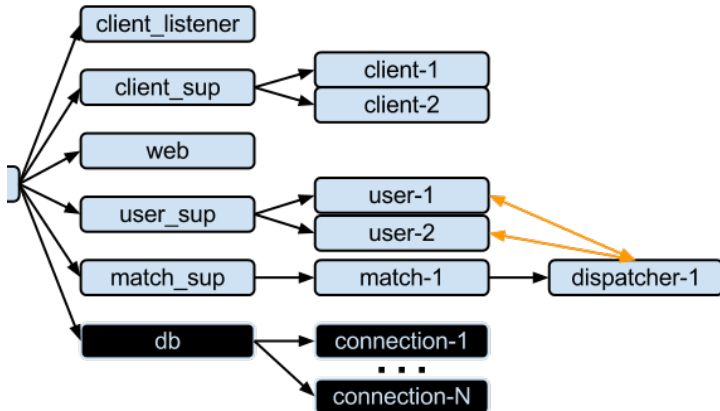
```
handle_call(Request, From, State) ->
  [RedisConn|Redis] = State#state.redis,
  proc_lib:spawn_link(
    fun() ->
      Res = handle_call(Request, RedisConn),
      gen_server:reply(From, Res)
    end),
  {noreply, State#state{redis =
    Redis ++ [RedisConn] }}.
```



## STAGE 2.1


### CONNECTION TWEAKS

#### SYSTEM ARCHITECTURE



# STAGE 2.1

## RESULTS



N 8000  
C 500  
ART 15s

## STAGE 2.2

GEN\_EVENT

SUP\_HANDLER

- Don't use it
- Monitor the processes instead

LONG DELIVERY QUEUES

- Use *repeaters*



## STAGE 2.2

GEN\_EVENT

SUP\_HANDLER

- Don't use it
- Monitor the processes instead

LONG DELIVERY QUEUES

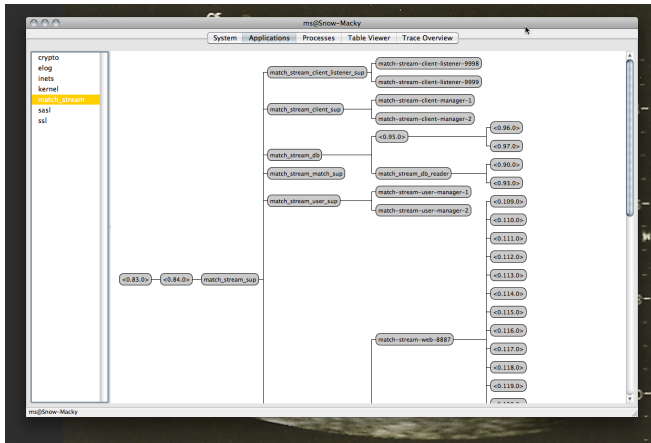
- Use *repeaters*



# STAGE 2.2

GEN\_EVENT


## SYSTEM ARCHITECTURE





## STAGE 2.2

### RESULTS



- **8192** users
- **256** at a time
- **8s** ART

## STAGE 2.3

### GEN\_SERVER

#### CALL TIMEOUTS

Remember `gen_server:reply/2`

#### MEMORY FOOTPRINT

Remember `hibernate`

#### LONG INIT/1

Use 0 timeout



## STAGE 2.3

### GEN\_SERVER

#### CALL TIMEOUTS

Remember `gen_server:reply/2`

#### MEMORY FOOTPRINT

Remember `hibernate`

#### LONG INIT/1

Use 0 timeout



## STAGE 2.3

### GEN\_SERVER

#### CALL TIMEOUTS

Remember `gen_server:reply/2`

#### MEMORY FOOTPRINT

Remember `hibernate`

#### LONG INIT/1

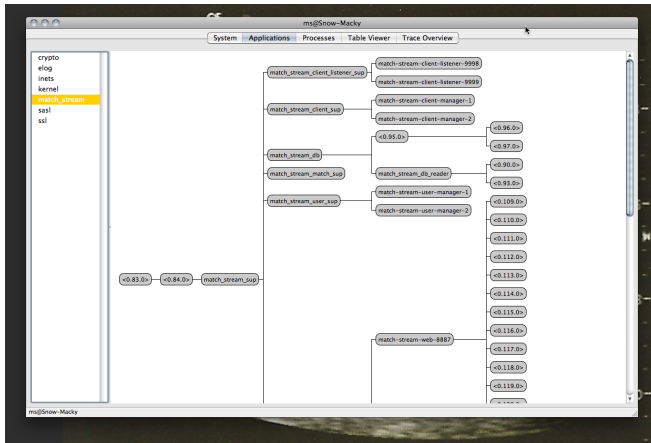
Use 0 timeout



# STAGE 2.3


GEN\_SERVER

## SYSTEM ARCHITECTURE



## STAGE 2.3

### RESULTS



- **32768** users
- **1024** at a time
- **1s** ART



## STAGE 2.4

### SUPERVISORS

- Sometimes `simple_one_for_one` supervisors get **overburdened** because they have too many children
- Try a supervisor hierarchy with several managers below the main supervisor
- Turn `supervisor:start_child/2` calls into something like

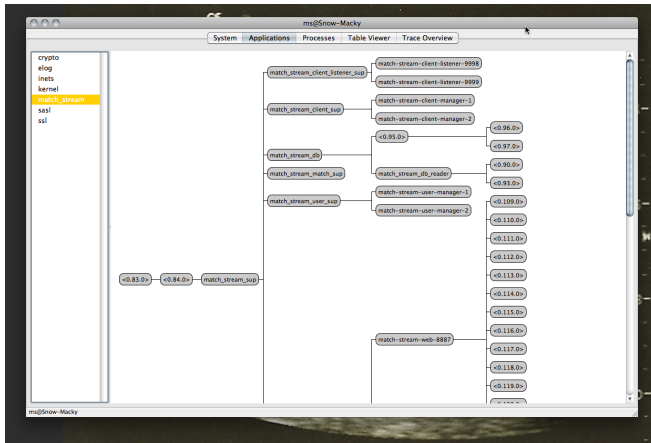
```
supervisor:start_child(  
    list_to_atom("module-name_" ++  
                                     integer_to_list(random:seed(1000000000)))
```



# STAGE 2.4

## SUPERVISORS


### SYSTEM ARCHITECTURE





## STAGE 2.4

### RESULTS



- **65536** users
- **2048** at a time
- **1s** ART



## STAGE 2.5

### OTHER PROCESSES

#### TIMERS

- Don't use the timer module
- Use `erlang:send_after`

#### LOGGING

- Don't log too much
- Use a good logging system

#### REGISTRATION

- Sometimes it's better to register processes instead of keeping track of their pids manually
- You can always register processes **both** locally and globally



## STAGE 2.5

### OTHER PROCESSES

#### TIMERS

- Don't use the `timer` module
- Use `erlang:send_after`

#### LOGGING

- Don't log too much
- Use a good logging system

#### REGISTRATION

- Sometimes it's better to register processes instead of keeping track of their pids manually
- You can always register processes **both** locally and globally



## STAGE 2.5

### OTHER PROCESSES

#### TIMERS

- Don't use the `timer` module
- Use `erlang:send_after`

#### LOGGING

- Don't log too much
- Use a good logging system

#### REGISTRATION

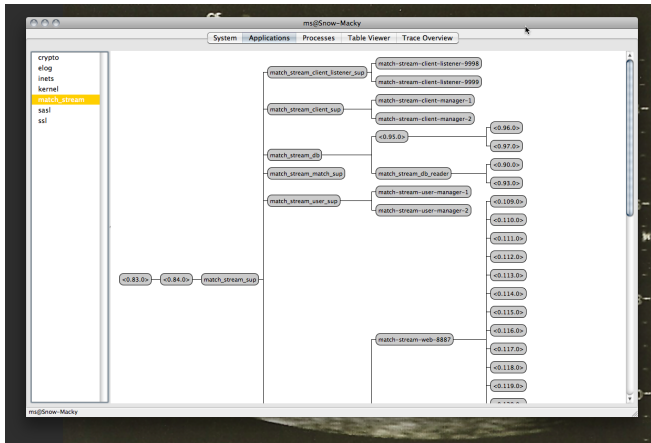
- Sometimes it's better to register processes instead of keeping track of their pids manually
- You can always register processes **both** locally and globally



## STAGE 2.5


### OTHER PROCESSES

## SYSTEM ARCHITECTURE



## STAGE 2.5

### RESULTS



- **65536** users
- **8192** at a time
- **10ms** ART

## STAGE 3

### ADDING NODES

#### GOALS

- Find the best system topology

#### STEPS

- Prepare the system to run in more than one node
- Decide if nodes should be connected or independent
- Decide if nodes should be on the same machine or not



## STAGE 3

### ADDING NODES

#### GOALS

- Find the best system topology

#### STEPS

- Prepare the system to run in more than one node
- Decide if nodes should be connected or independent
- Decide if nodes should be on the same machine or not

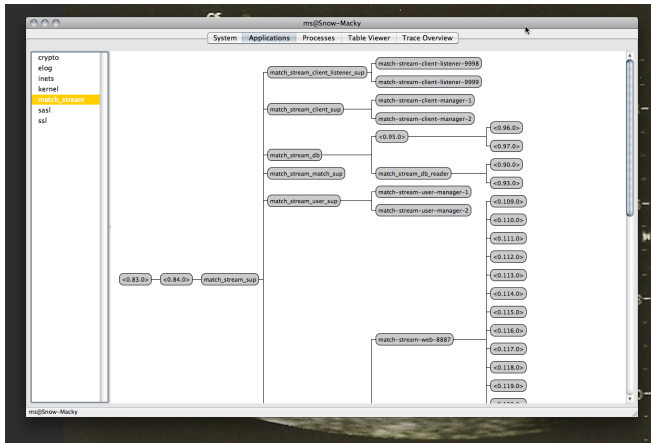




# STAGE 3


## ADDING NODES

### SYSTEM ARCHITECTURE



# STAGE 3

## RESULTS



- **100K** users
- **32768** at a time
- **10ms** ART



# SUMMARY

- This is an **iterative** process
- It worked awesomely for us in both experimental and real-life systems
- It's no **silver bullet**
- The list of *Tips and Tricks* grows **constantly** over time



# SUMMARY

- This is an **iterative** process
- It worked awesomely for us in both experimental and real-life systems
- It's no **silver bullet**
- The list of *Tips and Tricks* grows **constantly** over time



# SUMMARY

- This is an **iterative** process
- It worked awesomely for us in both experimental and real-life systems
- It's no **silver bullet**
- The list of *Tips and Tricks* grows **constantly** over time



# SUMMARY

- This is an **iterative** process
- It worked awesomely for us in both experimental and real-life systems
- It's no **silver bullet**
- The list of *Tips and Tricks* grows **constantly** over time



# SCALING TOPICS

THAT WEREN'T COVERED ON THIS PRESENTATION

- Managing many nodes
- Choosing databases
- System specific improvements
- Measuring tools



# QUESTIONS





# Thanks!

