

SCALING ERLANG WEB APPLICATIONS

100 TO 100K USERS AT ONE WEB SERVER

Fernando Benavides (*@elbrujohalcon*)

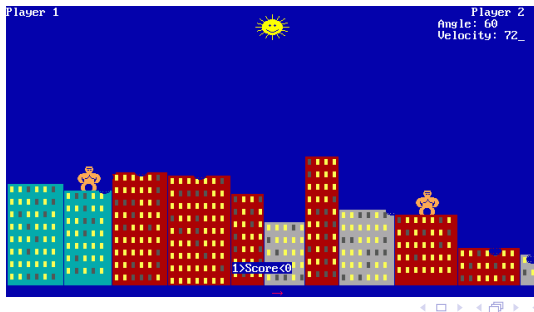
Inaka Labs

March 12, 2012



HELLO WORLD!

- I'm a developer since I was 10
- I'm an Erlang developer since 2008
- I've worked in many dynamic web sites
- Most of them with high scale requirements
- I'll share my experience with you



HELLO WORLD!

- I'm a developer since I was 10
- I'm an Erlang developer since 2008
- I've worked in many dynamic web sites
- Most of them with high scale requirements
- I'll share my experience with you



HELLO WORLD!

- I'm a developer since I was 10
- I'm an Erlang developer since 2008
- I've worked in many dynamic web sites
- Most of them with high scale requirements
- I'll share my experience with you



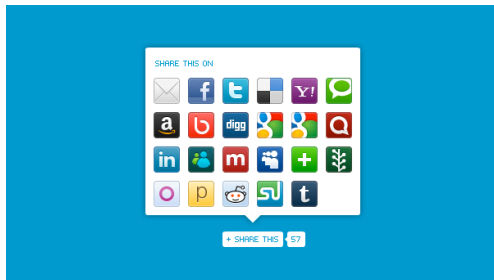
HELLO WORLD!

- I'm a developer since I was 10
- I'm an Erlang developer since 2008
- I've worked in many dynamic web sites
- Most of them with high scale requirements
- I'll share my experience with you



HELLO WORLD!

- I'm a developer since I was 10
- I'm an Erlang developer since 2008
- I've worked in many dynamic web sites
- Most of them with high scale requirements
- I'll share my experience with you



OUTLINE

THE CHALLENGE

What do we have to deal with?

THE PLAN

How do we face it?

THE TIPS AND TRICKS

What have we learned from it?



OUTLINE

THE CHALLENGE

What do we have to deal with?

THE PLAN

How do we face it?

THE TIPS AND TRICKS

What have we learned from it?



OUTLINE

THE CHALLENGE

What do we have to deal with?

THE PLAN

How do we face it?

THE TIPS AND TRICKS

What have we learned from it?



THE CHALLENGE

We will work on the scalability of a *web* project that has an *HTTP API* and keeps clients *connected* to the server for *long periods* of time.

Examples:

- Social sites
- Chat sites
- Sports sites



THE CHALLENGE

We will work on the scalability of a *web* project that has an *HTTP API* and keeps clients *connected* to the server for *long periods* of time.

Examples:

- Social sites
- Chat sites
- Sports sites



THE CHALLENGE

We will work on the scalability of a *web* project that has an *HTTP API* and keeps clients *connected* to the server *for long periods* of time.

Examples:

- Social sites
- Chat sites
- Sports sites



THE CHALLENGE

We will work on the scalability of a *web* project that has an *HTTP API* and keeps clients *connected* to the server for *long periods* of time.

Examples:

- Social sites
- Chat sites
- Sports sites



THE CHALLENGE

We will work on the scalability of a *web* project that has an *HTTP API* and keeps clients *connected* to the server for *long periods* of time.

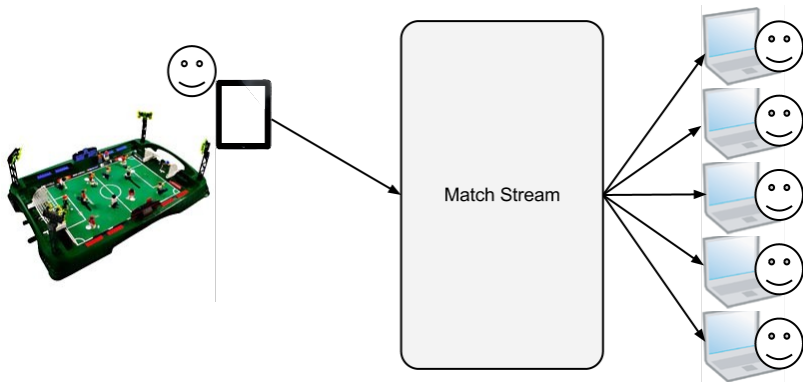
Examples:

- Social sites
- Chat sites
- Sports sites



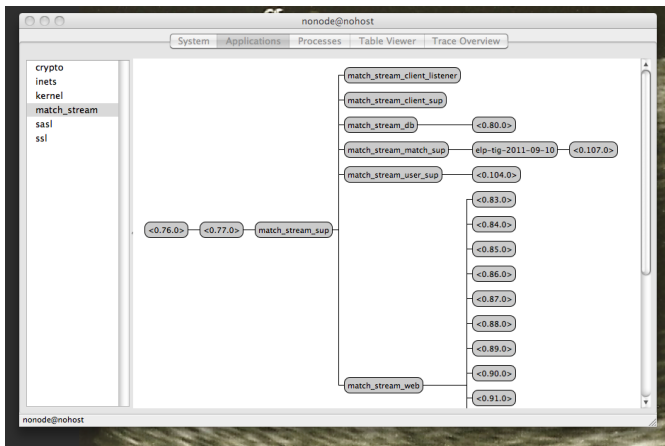
MATCH STREAM

THE SAMPLE PROJECT



MATCH STREAM

INITIAL ARCHITECTURE



THE CHALLENGE

We will focus on

- OTP behaviours
- TCP connections
- mochiweb
- Underlying system configurations

*We will **not** deal with*

- Multiple machines/nodes
- Databases



THE CHALLENGE

We will focus on

- OTP behaviours
- TCP connections
- mochiweb
- Underlying system configurations

*We will **not** deal with*

- Multiple machines/nodes
- Databases



THE PLAN



GENERAL CONSIDERATIONS

- Build a system that **works**
- Automate your clients
- Keep a human watching
- Be patient



GENERAL CONSIDERATIONS

- Build a system that **works**
- Automate your clients
- Keep a human watching
- Be patient



GENERAL CONSIDERATIONS

- Build a system that **works**
- Automate your clients
- **Keep a human watching**
- Be patient



GENERAL CONSIDERATIONS

- Build a system that **works**
- Automate your clients
- Keep a human watching
- Be patient



GENERAL CONSIDERATIONS

- Build a system that **works**
- Automate your clients
- Keep a human watching
- Be patient



STAGE 1

FINDING THE INITIAL BOUNDARIES

GOALS

- Test the system as it is
- How many users can the system handle **as is**?
- Find N and C



STAGE 1

FINDING THE INITIAL BOUNDARIES

STEPS

- Choose N and C
- Test the API
- Test long-lived connections
- Test both
- Repeat with higher values for N and C



STAGE 1

FINDING THE INITIAL BOUNDARIES

STEPS

- Choose N and C
- Test the API
- Test long-lived connections
- Test both
- Repeat with higher values for N and C



STAGE 2

BLACKBOX TESTS

GOALS

- Improve the system environment
- Tune-In the machine(s)
- **Don't** touch the code



STAGE 2

BLACKBOX TESTS

STEPS

- Check kernel variables
- Check system limits
- Check Erlang VM parameters
- Repeat Stage 1



STAGE 2

BLACKBOX TESTS

STEPS

- Check kernel variables
- Check system limits
- Check Erlang VM parameters
- Repeat Stage 1



STAGE 3

ERLANG TUNING

GOALS

- Tune up **your** system
- Discover scalability issues and fix them
- Find the biggest N and C for **one node**



STAGE 3

ERLANG TUNING

STEPS

- Choose N and C to fail
- Find a problem
- Fix it
- Add it to the list of *Tips and Tricks*
- Repeat with higher values for N and C



STAGE 3

ERLANG TUNING

STEPS

- Choose N and C to fail
- Find a problem
- Fix it
- Add it to the list of *Tips and Tricks*
- Repeat with higher values for N and C



STAGE 4

ADDING NODES

GOALS

- Get the system ready to work on many nodes
- Design the system topology
- Find N and C **per node**



STAGE 4

ADDING NODES

STEPS

- Get the second node running
- Choose N and C
- Try interconnected instances
- Try independent instances
- Repeat with higher values for N and C



STAGE 4

ADDING NODES

STEPS

- Get the second node running
- Choose N and C
- Try interconnected instances
- Try independent instances
- Repeat with higher values for N and C



TIPS AND TRICKS



OS TWEAKS

Kernel Variables

```
sysctl -w net.ipv4.ip_local_port_range="1024 65535"  
sysctl -w net.core.rmem_max=16777216  
sysctl -w net.core.wmem_max=16777216  
sysctl -w net.ipv4.tcp_rmem="4096 87380 16777216"  
sysctl -w net.ipv4.tcp_wmem="4096 65536 16777216"  
sysctl -w net.ipv4.tcp_syncookies=1  
sysctl -w net.ipv4.tcp_mem="50576 64768 98152"  
sysctl -w net.core.netdev_max_backlog=2500  
sysctl -w net.netfilter.nf_conntrack_max=1233000
```

Open Files Limit

```
ulimit -n 999999
```

Erlang VM tweaks

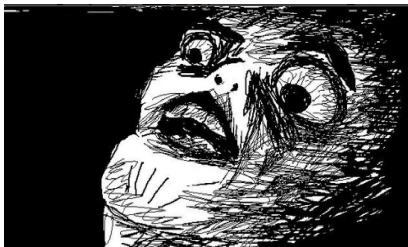
+P Number of Processes
+K Kernell Polling
-SMP SMP Support



OS TWEAKS

Kernel Variables

```
sysctl -w net.ipv4.ip_local_port_range="1024 65535"  
sysctl -w net.core.rmem_max=16777216  
sysctl -w net.core.wmem_max=16777216  
sysctl -w net.ipv4.tcp_rmem="4096 87380 16777216"  
sysctl -w net.ipv4.tcp_wmem="4096 65536 16777216"  
sysctl -w net.ipv4.tcp_synccookies=1  
sysctl -w net.ipv4.tcp_mem="50576 64768 98152"  
sysctl -w net.core.netdev_max_backlog=2500  
sysctl -w net.netfilter.nf_conntrack_max=1233000
```



OS TWEAKS

Kernel Variables

```
sysctl -w net.ipv4.ip_local_port_range="1024 65535"  
sysctl -w net.core.rmem_max=16777216  
sysctl -w net.core.wmem_max=16777216  
sysctl -w net.ipv4.tcp_rmem="4096 87380 16777216"  
sysctl -w net.ipv4.tcp_wmem="4096 65536 16777216"  
sysctl -w net.ipv4.tcp_syncookies=1  
sysctl -w net.ipv4.tcp_mem="50576 64768 98152"  
sysctl -w net.core.netdev_max_backlog=2500  
sysctl -w net.netfilter.nf_conntrack_max=1233000
```

Open Files Limit

```
ulimit -n 999999
```

Erlang VM tweaks

- +P Number of Processes
- +K Kernell Polling
- SMP SMP Support



OS TWEAKS

Kernel Variables

```
sysctl -w net.ipv4.ip_local_port_range="1024 65535"  
sysctl -w net.core.rmem_max=16777216  
sysctl -w net.core.wmem_max=16777216  
sysctl -w net.ipv4.tcp_rmem="4096 87380 16777216"  
sysctl -w net.ipv4.tcp_wmem="4096 65536 16777216"  
sysctl -w net.ipv4.tcp_syncookies=1  
sysctl -w net.ipv4.tcp_mem="50576 64768 98152"  
sysctl -w net.core.netdev_max_backlog=2500  
sysctl -w net.netfilter.nf_conntrack_max=1233000
```

Open Files Limit

```
ulimit -n 999999
```

Erlang VM tweaks

- +P Number of Processes
- +K Kernell Polling
- SMP SMP Support



CONNECTION TWEAKS

BACKLOG

- Allow more concurrent connections
- Remember HTTP *runs on* TCP

CONNECTIONS

- Don't use just one of them
- Check inbound and outbound connections



CONNECTION TWEAKS

BACKLOG

- Allow more concurrent connections
- Remember HTTP *runs on* TCP

CONNECTIONS

- Don't use just one of them
- Check inbound and outbound connections



GEN_EVENT

SUP_HANDLER

- Don't use it
- Monitor the processes instead

LONG DELIVERY QUEUES

- Use *repeaters*



GEN_EVENT

SUP_HANDLER

- Don't use it
- Monitor the processes instead

LONG DELIVERY QUEUES

- Use *repeaters*



GEN_SERVER

CALL TIMEOUTS

Remember `gen_server:reply/2`

MEMORY FOOTPRINT

Remember `hibernate`

LONG INIT/1

Use 0 timeout



GEN_SERVER

CALL TIMEOUTS

Remember `gen_server:reply/2`

MEMORY FOOTPRINT

Remember `hibernate`

LONG INIT/1

Use 0 timeout



GEN_SERVER

CALL TIMEOUTS

Remember `gen_server:reply/2`

MEMORY FOOTPRINT

Remember `hibernate`

LONG INIT/1

Use 0 timeout



SUPERVISORS

- Sometimes `simple_one_for_one` supervisors get **overburdened** because they have too many children
- Try a supervisor hierarchy with several managers below the main supervisor
- Turn `supervisor:start_child/2` calls into something like

```
supervisor:start_child(  
  list_to_atom("module-name_" ++  
              integer_to_list(random:uniform(#ofSupervisors))).
```



OTHER PROCESSES

TIMERS

- Don't use the timer module
- Use `erlang:send_after`

LOGGING

- Don't log too much
- Use a good logging system

REGISTRATION

- Sometimes it's better to register processes instead of keeping track of their pids manually
- You can always register processes **both** locally and globally



OTHER PROCESSES

TIMERS

- Don't use the timer module
- Use `erlang:send_after`

LOGGING

- Don't log too much
- Use a good logging system

REGISTRATION

- Sometimes it's better to register processes instead of keeping track of their pids manually
- You can always register processes **both** locally and globally



OTHER PROCESSES

TIMERS

- Don't use the `timer` module
- Use `erlang:send_after`

LOGGING

- Don't log too much
- Use a good logging system

REGISTRATION

- Sometimes it's better to register processes instead of keeping track of their pids manually
- You can always register processes **both** locally and globally



SUMMARY

- It is an **iterative** process
- It worked awesomely for us in both experimental and real-life systems
- It's no **silver bullet**
- The list of *Tips and Tricks* grows **constantly** over time



SUMMARY

- It is an **iterative** process
- It worked awesomely for us in both experimental and real-life systems
- It's no **silver bullet**
- The list of *Tips and Tricks* grows **constantly** over time



SUMMARY

- It is an **iterative** process
- It worked awesomely for us in both experimental and real-life systems
- It's no **silver bullet**
- The list of *Tips and Tricks* grows **constantly** over time



SUMMARY

- It is an **iterative** process
- It worked awesomely for us in both experimental and real-life systems
- It's no **silver bullet**
- The list of *Tips and Tricks* grows **constantly** over time



SCALING TOPICS

THAT WEREN'T COVERED ON THIS PRESENTATION

- Adding nodes
- Choosing databases
- System specific improvements
- Measuring tools



QUESTIONS



Thanks!

