

SCALING ERLANG WEB APPLICATIONS

100 TO 100K USERS AT ONE WEB SERVER

Fernando Benavides (*@elbrujohalcon*)

Inaka Labs

March 20, 2012



INAKA NETWORKS

presents ...



INAKA NETWORKS

presents ...



El Brujo Halcón

in ...



El Brujo Halcón

in ...



SCALING ERLANG

Based on a true story



SCALING ERLANG

Based on a true story



A not so long time ago



A not so long time ago



in a country far far away. . .



A FRIEND

Hey! Let's watch the *superclásico*!!!

BRUJO

I can't, I'm at *the office*

A FRIEND

...

BRUJO

We need an app for that!



A FRIEND

Hey! Let's watch the *superclásico*!!!

BRUJO

I can't, I'm at *the office*

A FRIEND

...

BRUJO

We need an app for that!



A FRIEND

Hey! Let's watch the *superclásico*!!!

BRUJO

I can't, I'm at *the office*

A FRIEND

...

BRUJO

We need an app for that!



A FRIEND

Hey! Let's watch the *superclásico*!!!

BRUJO

I can't, I'm at *the office*

A FRIEND

...

BRUJO

We need an app for that!



BRUJO

Let's call it MATCHSTREAM

A FRIEND

Ok, then. . . We know there will be hundreds of thousands of users, right?

We need the system to **scale**

BRUJO

Of course! We should use Erlang!



BRUJO

Let's call it MATCHSTREAM

A FRIEND

Ok, then... We know there will be hundreds of thousands of users, right?

We need the system to **scale**

BRUJO

Of course! We should use Erlang!



BRUJO

Let's call it MATCHSTREAM

A FRIEND

Ok, then... We know there will be hundreds of thousands of users, right?

We need the system to **scale**

BRUJO

Of course! We should use Erlang!




A while later...



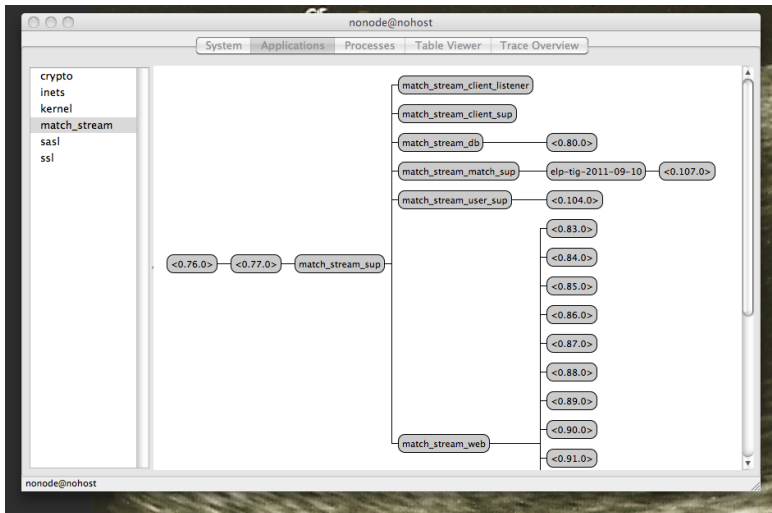
MATCHSTREAM

SYSTEM DESCRIPTION



MATCHSTREAM

ARCHITECTURE



TODO: Take this picture with client(s) connected



BRUJO

Boca plays again today, let's try our system out with this game!

What can **possibly** go wrong?

USER 1

Wow! MATCHSTREAM is awesome!

...

USER 100

Hey! this system is a total crap! It doesn't even let me connect to it!

BRUJO

WTF?! The system doesn't scale!!

A FRIEND

Didn't you use Erlang?



BRUJO

Boca plays again today, let's try our system out with this game!

What can **possibly** go wrong?

USER 1

Wow! MATCHSTREAM is awesome!

...

USER 100

Hey! this system is a total crap! It doesn't even let me connect to it!

BRUJO

WTF?! The system doesn't scale!!

A FRIEND

Didn't you use Erlang?



BRUJO

Boca plays again today, let's try our system out with this game!

What can **possibly** go wrong?

USER 1

Wow! MATCHSTREAM is awesome!

...

USER 100

Hey! this system is a total crap! It doesn't even let me connect to it!

BRUJO

WTF?! The system doesn't scale!!

A FRIEND

Didn't you use Erlang?



BRUJO

Boca plays again today, let's try our system out with this game!

What can **possibly** go wrong?

USER 1

Wow! MATCHSTREAM is awesome!

...

USER 100

Hey! this system is a total crap! It doesn't even let me connect to it!

BRUJO

WTF?! The system doesn't scale!!

A FRIEND

Didn't you use Erlang?



BRUJO

Boca plays again today, let's try our system out with this game!

What can **possibly** go wrong?

USER 1

Wow! MATCHSTREAM is awesome!

...

USER 100

Hey! this system is a total crap! It doesn't even let me connect to it!

BRUJO

WTF?! **The system doesn't scale!!**

A FRIEND

Didn't you use Erlang?



BRUJO

Boca plays again today, let's try our system out with this game!

What can **possibly** go wrong?

USER 1

Wow! MATCHSTREAM is awesome!

...

USER 100

Hey! this system is a total crap! It doesn't even let me connect to it!

BRUJO

WTF?! **The system doesn't scale!!**

A FRIEND

Didn't you use **Erlang**?



LESSON LEARNED

Just using Erlang is not enough to make your system scale



So, we made it scale...



We made sure the system was working.

- We built a simulator
- We improved the logging mechanisms
- We tested the system



We made sure the system was working.

- We built a simulator
- We improved the logging mechanisms
- We tested the system



We made sure the system was working.

- We built a simulator
- We improved the logging mechanisms
- We tested the system



We made sure the system was working.

- We built a simulator
- We improved the logging mechanisms
- We tested the system



1024 users / 4 at a time



The system is fine, let's tune up the server where it's installed

So, we checked the kernel variables and system limits for

- Concurrent TCP connections
- Open files limit
- TCP backlog size
- Swap memory allocation
- Erlang VM process limit



The system is fine, let's tune up the server where it's installed
So, we checked the kernel variables and system limits for

- Concurrent TCP connections
- Open files limit
- TCP backlog size
- TCP memory allocation
- Erlang VM process limit



The system is fine, let's tune up the server where it's installed
So, we checked the kernel variables and system limits for

- Concurrent TCP connections
- Open files limit
- TCP backlog size
- TCP memory allocation
- Erlang VM process limit



The system is fine, let's tune up the server where it's installed
So, we checked the kernel variables and system limits for

- Concurrent TCP connections
- Open files limit
- TCP backlog size
- TCP memory allocation
- Erlang VM process limit



The system is fine, let's tune up the server where it's installed
So, we checked the kernel variables and system limits for

- Concurrent TCP connections
- Open files limit
- TCP backlog size
- TCP memory allocation
- Erlang VM process limit



The system is fine, let's tune up the server where it's installed
So, we checked the kernel variables and system limits for

- Concurrent TCP connections
- Open files limit
- TCP backlog size
- TCP memory allocation
- Erlang VM process limit



4096 users / 4 at a time



I've got a friend that may help us, he has a bag with several tips and tricks for us... MacGyver



I've got a friend that may help us, he has a bag with several tips and tricks for us... MacGyver



STEP 3

CONNECTION TWEAKS

BACKLOG

- Allow more concurrent connections
- Remember HTTP *runs on* TCP

CONNECTIONS

- Don't use just one of them
- Check inbound and outbound connections



STEP 3

CONNECTION TWEAKS

BACKLOG

- Allow more concurrent connections
- Remember HTTP *runs on* TCP

CONNECTIONS

- Don't use just one of them
- Check inbound and outbound connections



TODO users / TODO at a time



SUP_HANDLER

- Don't use it
- Monitor the processes instead

LONG DELIVERY QUEUES

- Use *repeaters*



SUP_HANDLER

- Don't use it
- Monitor the processes instead

LONG DELIVERY QUEUES

- Use *repeaters*



TODO users / TODO at a time



CALL TIMEOUTS

Remember `gen_server:reply/2`

MEMORY FOOTPRINT

Remember `hibernate`

LONG INIT/1

Use 0 timeout



CALL TIMEOUTS

Remember `gen_server:reply/2`

MEMORY FOOTPRINT

Remember `hibernate`

LONG INIT/1

Use 0 timeout



CALL TIMEOUTS

Remember `gen_server:reply/2`

MEMORY FOOTPRINT

Remember `hibernate`

LONG INIT/1

Use 0 timeout



TODO users / TODO at a time



- Sometimes `simple_one_for_one` supervisors get **overburdened** because they have too many children
- Try a supervisor hierarchy with several managers below the main supervisor
- Turn `supervisor:start_child/2` calls into something like

```
supervisor:start_child(  
    list_to_atom("module-name_" ++  
                integer_to_list(random:uniform(#ofSupervisors))).
```



TODO users / TODO at a time



TIMERS

- Don't use the `timer` module
- Use `erlang:send_after`

LOGGING

- Don't log too much
- Use a good logging system

REGISTRATION

- Sometimes it's better to register processes instead of keeping track of their pids manually
- You can always register processes **both** locally and globally



TIMERS

- Don't use the `timer` module
- Use `erlang:send_after`

LOGGING

- Don't log too much
- Use a good logging system

REGISTRATION

- Sometimes it's better to register processes instead of keeping track of their pids manually
- You can always register processes **both** locally and globally



TIMERS

- Don't use the `timer` module
- Use `erlang:send_after`

LOGGING

- Don't log too much
- Use a good logging system

REGISTRATION

- Sometimes it's better to register processes instead of keeping track of their pids manually
- You can always register processes **both** locally and globally



64000 users / 8000 at a time



TODO: Img of what the system looks like at this point



STEP 4

Well, let's add some nodes to it!



STEP 4

ADDING NODES

Again, it's not as easy as just starting the app in another Erlang node We needed to find the best topology, we considered using:

- connected nodes
- independent nodes

We had to decide which processes needed to communicate and how and of course, test the whole system again



STEP 4

ADDING NODES

Again, it's not as easy as just starting the app in another Erlang node. We needed to find the best topology, we considered using:

- connected nodes
- independent nodes

We had to decide which processes needed to communicate and how and of course, test the whole system again



STEP 4

ADDING NODES

Again, it's not as easy as just starting the app in another Erlang node. We needed to find the best topology, we considered using:

- connected nodes
- independent nodes

We had to decide which processes needed to communicate and how and of course, test the whole system again



STEP 4

ADDING NODES

Again, it's not as easy as just starting the app in another Erlang node. We needed to find the best topology, we considered using:

- connected nodes
- independent nodes

We had to decide which processes needed to communicate and how and of course, test the whole system again



25000 users per node / 8000 per computer at a time with 4
nodes on the same computer... 100K users / 8000 at a time



25000 users per node / 8000 per computer at a time with 4
nodes on the same computer... 100K users / 8000 at a time



25000 users per node / 8000 per computer at a time with 4
nodes on the same computer... 100K users / 8000 at a time

