

SCALING ERLANG WEB APPLICATIONS

100 TO 100K USERS AT ONE WEB SERVER

Fernando Benavides (*@elbrujohalcon*)

Inaka Labs

March 20, 2012



INAKA NETWORKS

presents ...



INAKA NETWORKS

presents ...



El Brujo Halcón

in ...



El Brujo Halcón

in ...



SCALING ERLANG

Based on a true story



SCALING ERLANG

Based on a true story



A not so long time ago



A *not* so long time ago



in a country far far away. . .



A FRIEND

Hey! Let's watch the *superclásico*!!!

BRUJO

I can't, I'm at *the office*

A FRIEND

...

BRUJO

We need an app for that!



A FRIEND

Hey! Let's watch the *superclásico*!!!

BRUJO

I can't, I'm at *the office*

A FRIEND

...

BRUJO

We need an app for that!



A FRIEND

Hey! Let's watch the *superclásico*!!!

BRUJO

I can't, I'm at *the office*

A FRIEND

...

BRUJO

We need an app for that!



A FRIEND

Hey! Let's watch the *superclásico*!!!

BRUJO

I can't, I'm at *the office*

A FRIEND

...

BRUJO

We need an app for that!



BRUJO

Let's call it MATCHSTREAM

A FRIEND

Ok, then. . . We know there will be hundreds of thousands of users, right?

We need the system to **scale**

BRUJO

Of course! We should use Erlang!



BRUJO

Let's call it MATCHSTREAM

A FRIEND

Ok, then... We know there will be hundreds of thousands of users, right?

We need the system to **scale**

BRUJO

Of course! We should use Erlang!



BRUJO

Let's call it MATCHSTREAM

A FRIEND

Ok, then... We know there will be hundreds of thousands of users, right?

We need the system to **scale**

BRUJO

Of course! We should use Erlang!




A while later...



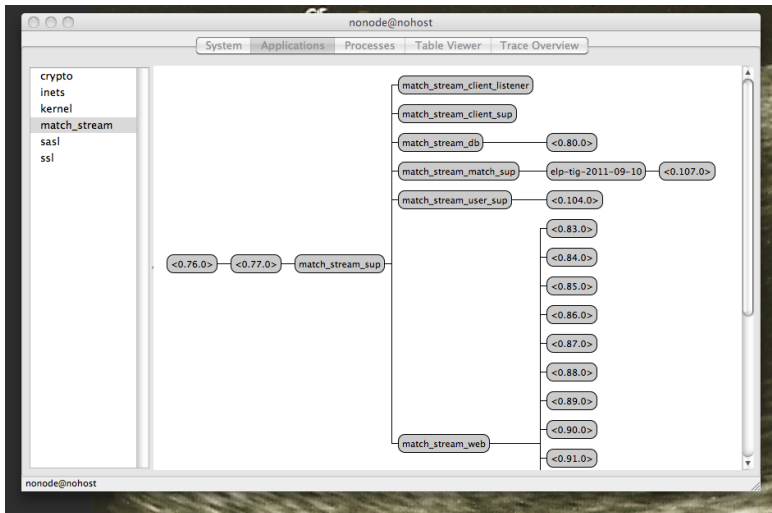
MATCHSTREAM

SYSTEM DESCRIPTION



MATCHSTREAM

ARCHITECTURE



TODO: Take this picture with client(s) connected



BRUJO

Boca plays again today, let's try our system out with this game!

What can **possibly** go wrong?

USER 1

Wow! MATCHSTREAM is awesome!

...

USER 100

Hey! this system is a total crap! It doesn't even let me connect to it!

BRUJO

WTF?! The system doesn't scale!!

A FRIEND

Didn't you use Erlang?



BRUJO

Boca plays again today, let's try our system out with this game!

What can **possibly** go wrong?

USER 1

Wow! MATCHSTREAM is awesome!

...

USER 100

Hey! this system is a total crap! It doesn't even let me connect to it!

BRUJO

WTF?! The system doesn't scale!!

A FRIEND

Didn't you use Erlang?



BRUJO

Boca plays again today, let's try our system out with this game!

What can **possibly** go wrong?

USER 1

Wow! MATCHSTREAM is awesome!

...

USER 100

Hey! this system is a total crap! It doesn't even let me connect to it!

BRUJO

WTF?! The system doesn't scale!!

A FRIEND

Didn't you use Erlang?



BRUJO

Boca plays again today, let's try our system out with this game!

What can **possibly** go wrong?

USER 1

Wow! MATCHSTREAM is awesome!

...

USER 100

Hey! this system is a total crap! It doesn't even let me connect to it!

BRUJO

WTF?! The system doesn't scale!!

A FRIEND

Didn't you use Erlang?



BRUJO

Boca plays again today, let's try our system out with this game!

What can **possibly** go wrong?

USER 1

Wow! MATCHSTREAM is awesome!

...

USER 100

Hey! this system is a total crap! It doesn't even let me connect to it!

BRUJO

WTF?! **The system doesn't scale!!**

A FRIEND

Didn't you use Erlang?



BRUJO

Boca plays again today, let's try our system out with this game!

What can **possibly** go wrong?

USER 1

Wow! MATCHSTREAM is awesome!

...

USER 100

Hey! this system is a total crap! It doesn't even let me connect to it!

BRUJO

WTF?! **The system doesn't scale!!**

A FRIEND

Didn't you use **Erlang**?



LESSON LEARNED

Just using Erlang is not enough to make your system scale



So, we made it scale...



First of all we wanted to be sure that the system was actually working.

- We built a simulator
- We improved the logging mechanisms
- We tested the system
- We found its initial scale limits



First of all we wanted to be sure that the system was actually working.

- We built a simulator
- We improved the logging mechanisms
- We tested the system
- We found its initial scale limits



First of all we wanted to be sure that the system was actually working.

- We built a simulator
- We improved the logging mechanisms
- We tested the system
- We found its initial scale limits



First of all we wanted to be sure that the system was actually working.

- We built a simulator
- We improved the logging mechanisms
- We tested the system
- We found its initial scale limits



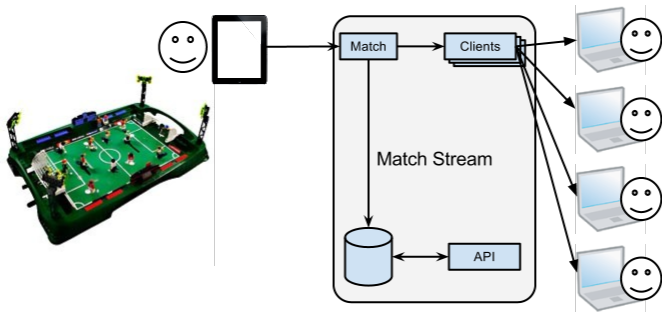
First of all we wanted to be sure that the system was actually working.

- We built a simulator
- We improved the logging mechanisms
- We tested the system
- We found its initial scale limits



STEP 1

RESULTS



$N = 1024 / C = 4$



Once we knew the system was fine, we decided to tune up the server where it was installed. So, we checked the kernel variables and system limits for

- Concurrent TCP connections
- Open files limit
- TCP backlog size
- TCP memory allocation
- Erlang VM processes



Once we knew the system was fine, we decided to tune up the server where it was installed. So, we checked the kernel variables and system limits for

- Concurrent TCP connections
- Open files limit
- TCP backlog size
- TCP memory allocation
- Erlang VM process limit



Once we knew the system was fine, we decided to tune up the server where it was installed. So, we checked the kernel variables and system limits for

- Concurrent TCP connections
- Open files limit
- TCP backlog size
- TCP memory allocation
- Erlang VM process limit



Once we knew the system was fine, we decided to tune up the server where it was installed. So, we checked the kernel variables and system limits for

- Concurrent TCP connections
- Open files limit
- TCP backlog size
- TCP memory allocation
- Erlang VM process limit



Once we knew the system was fine, we decided to tune up the server where it was installed. So, we checked the kernel variables and system limits for

- Concurrent TCP connections
- Open files limit
- TCP backlog size
- TCP memory allocation
- Erlang VM process limit




Once we knew the system was fine, we decided to tune up the server where it was installed. So, we checked the kernel variables and system limits for

- Concurrent TCP connections
- Open files limit
- TCP backlog size
- TCP memory allocation
- Erlang VM process limit



STEP 2

RESULTS



$$N = 4096 / C = 4$$



STEP 3

Then we decided to start improving the different components of the system.

We called a friend to help us. . .



STEP 3

Then we decided to start improving the different components of the system.

We called a friend to help us...



STEP 3

Then we decided to start improving the different components of the system.

We called a friend to help us...



STEP 3

CONNECTION TWEAKS

BACKLOG

- Allow more concurrent connections
- Remember HTTP *runs on* TCP

CONNECTIONS

- Don't use just one of them
- Check inbound and outbound connections



STEP 3

CONNECTION TWEAKS

BACKLOG

- Allow more concurrent connections
- Remember HTTP *runs on* TCP

CONNECTIONS

- Don't use just one of them
- Check inbound and outbound connections



TODO users / TODO at a time



SUP_HANDLER

- Don't use it
- Monitor the processes instead

LONG DELIVERY QUEUES

- Use *repeaters*



SUP_HANDLER

- Don't use it
- Monitor the processes instead

LONG DELIVERY QUEUES

- Use *repeaters*



TODO users / TODO at a time



CALL TIMEOUTS

Remember `gen_server:reply/2`

MEMORY FOOTPRINT

Remember `hibernate`

LONG INIT/1

Use 0 timeout



CALL TIMEOUTS

Remember `gen_server:reply/2`

MEMORY FOOTPRINT

Remember `hibernate`

LONG INIT/1

Use 0 timeout



CALL TIMEOUTS

Remember `gen_server:reply/2`

MEMORY FOOTPRINT

Remember `hibernate`

LONG INIT/1

Use 0 timeout



TODO users / TODO at a time



- Sometimes `simple_one_for_one` supervisors get **overburdened** because they have too many children
- Try a supervisor hierarchy with several managers below the main supervisor
- Turn `supervisor:start_child/2` calls into something like

```
supervisor:start_child(  
  list_to_atom("module-name_" ++  
               integer_to_list(random:uniform(#ofSupervisors))).
```



TODO users / TODO at a time



TIMERS

- Don't use the `timer` module
- Use `erlang:send_after`

LOGGING

- Don't log too much
- Use a good logging system

REGISTRATION

- Sometimes it's better to register processes instead of keeping track of their pids manually
- You can always register processes **both** locally and globally



TIMERS

- Don't use the `timer` module
- Use `erlang:send_after`

LOGGING

- Don't log too much
- Use a good logging system

REGISTRATION

- Sometimes it's better to register processes instead of keeping track of their pids manually
- You can always register processes **both** locally and globally



TIMERS

- Don't use the `timer` module
- Use `erlang:send_after`

LOGGING

- Don't log too much
- Use a good logging system


REGISTRATION

- Sometimes it's better to register processes instead of keeping track of their pids manually
- You can always register processes **both** locally and globally



STEP 3

RESULTS



$N = 65536 / C = 8192$



TODO: Img of what the system looks like at this point



Well, let's add some nodes to it!



STEP 4

ADDING NODES

Again, it's not as easy as just starting the app in another Erlang node We needed to find the best topology, we considered using:

- connected nodes
- independent nodes

We had to decide which processes needed to communicate and how and of course, test the whole system again



STEP 4

ADDING NODES

Again, it's not as easy as just starting the app in another Erlang node. We needed to find the best topology, we considered using:

- connected nodes
- independent nodes

We had to decide which processes needed to communicate and how and of course, test the whole system again



STEP 4

ADDING NODES

Again, it's not as easy as just starting the app in another Erlang node. We needed to find the best topology, we considered using:

- connected nodes
- independent nodes

We had to decide which processes needed to communicate and how and of course, test the whole system again



STEP 4

ADDING NODES

Again, it's not as easy as just starting the app in another Erlang node We needed to find the best topology, we considered using:

- connected nodes
- independent nodes

We had to decide which processes needed to communicate and how and of course, test the whole system again



